

# Projeto - Processamento Digital de Imagens

Alunos: Gabriel de Oliveira e Silva Padovani Barreto RA: 2326221 Johanna Kirchner RA: 2265010

[https://github.com/Barreto-G/image\\_classifier\\_pdi](https://github.com/Barreto-G/image_classifier_pdi)

## Parte 1: Definição do tema, coleta e pré-processamento de imagens

Foram escolhidos itens de armarinho como tema dos objetos que formam o dataset. Foram determinadas 10 classes de objetos, cada uma contendo 5 itens dos quais foram capturadas 4 fotos utilizando uma camera digital, com fundo e posicionamento variados, combinando dois tipos de fundo com duas posições diferentes. As imagens foram armazenadas em pastas de acordo com sua classe, e nomeadas seguindo a formatação <CLASSID>-<IMG\_SEQUENCE>-V1|V2-B|W.png .

A exibição das imagens e os metadados do projeto podem ser visualizados no arquivo `show_dataset_info.ipynb` .


```
In [1]: from utilidades import *  
from trim_dataset import TrimDataset
```

```
In [2]: dataset = TrimDataset('trims_dataset')  
plot_class_grid(dataset, 2, 'nobbbox')
```



Para a marcação das imagens, utilizou-se o CVAT, onde foi marcado o objeto de interesse em cada uma das imagens.

```
In [3]: plot_class_grid(dataset, 2, 'bbox')
```



Para realizar o pré-processamento das imagens, e a subsequente criação do dataset, criou-se a classe `TrimDataset` , que armazena todas as informações relevantes do

Dataset e os métodos utilizados para o seu pré-processamento. Tal classe está localizada no arquivo `trims_dataset.py`. A primeira etapa do pré-processamento foi diminuir a resolução das imagens capturadas, tanto para diminuir o espaço ocupado pelos arquivos quando para acelerar o treinamento do modelo. A resolução das imagens originais foi diminuída de 2112x2112 para 224x224 sem perdas de eficiência do modelo.


Essa primeira etapa é realizada pela função `resize_image`, aplicada já na instancia de um objeto `TrimDataset`. Para reduzir o gasto computacional e de memória, as imagens incluídas nos arquivos do projeto já estão com a resolução reduzida, podendo ser encontradas na pasta `trims_dataset` ou arquivo `.tar.gz` de mesmo nome.

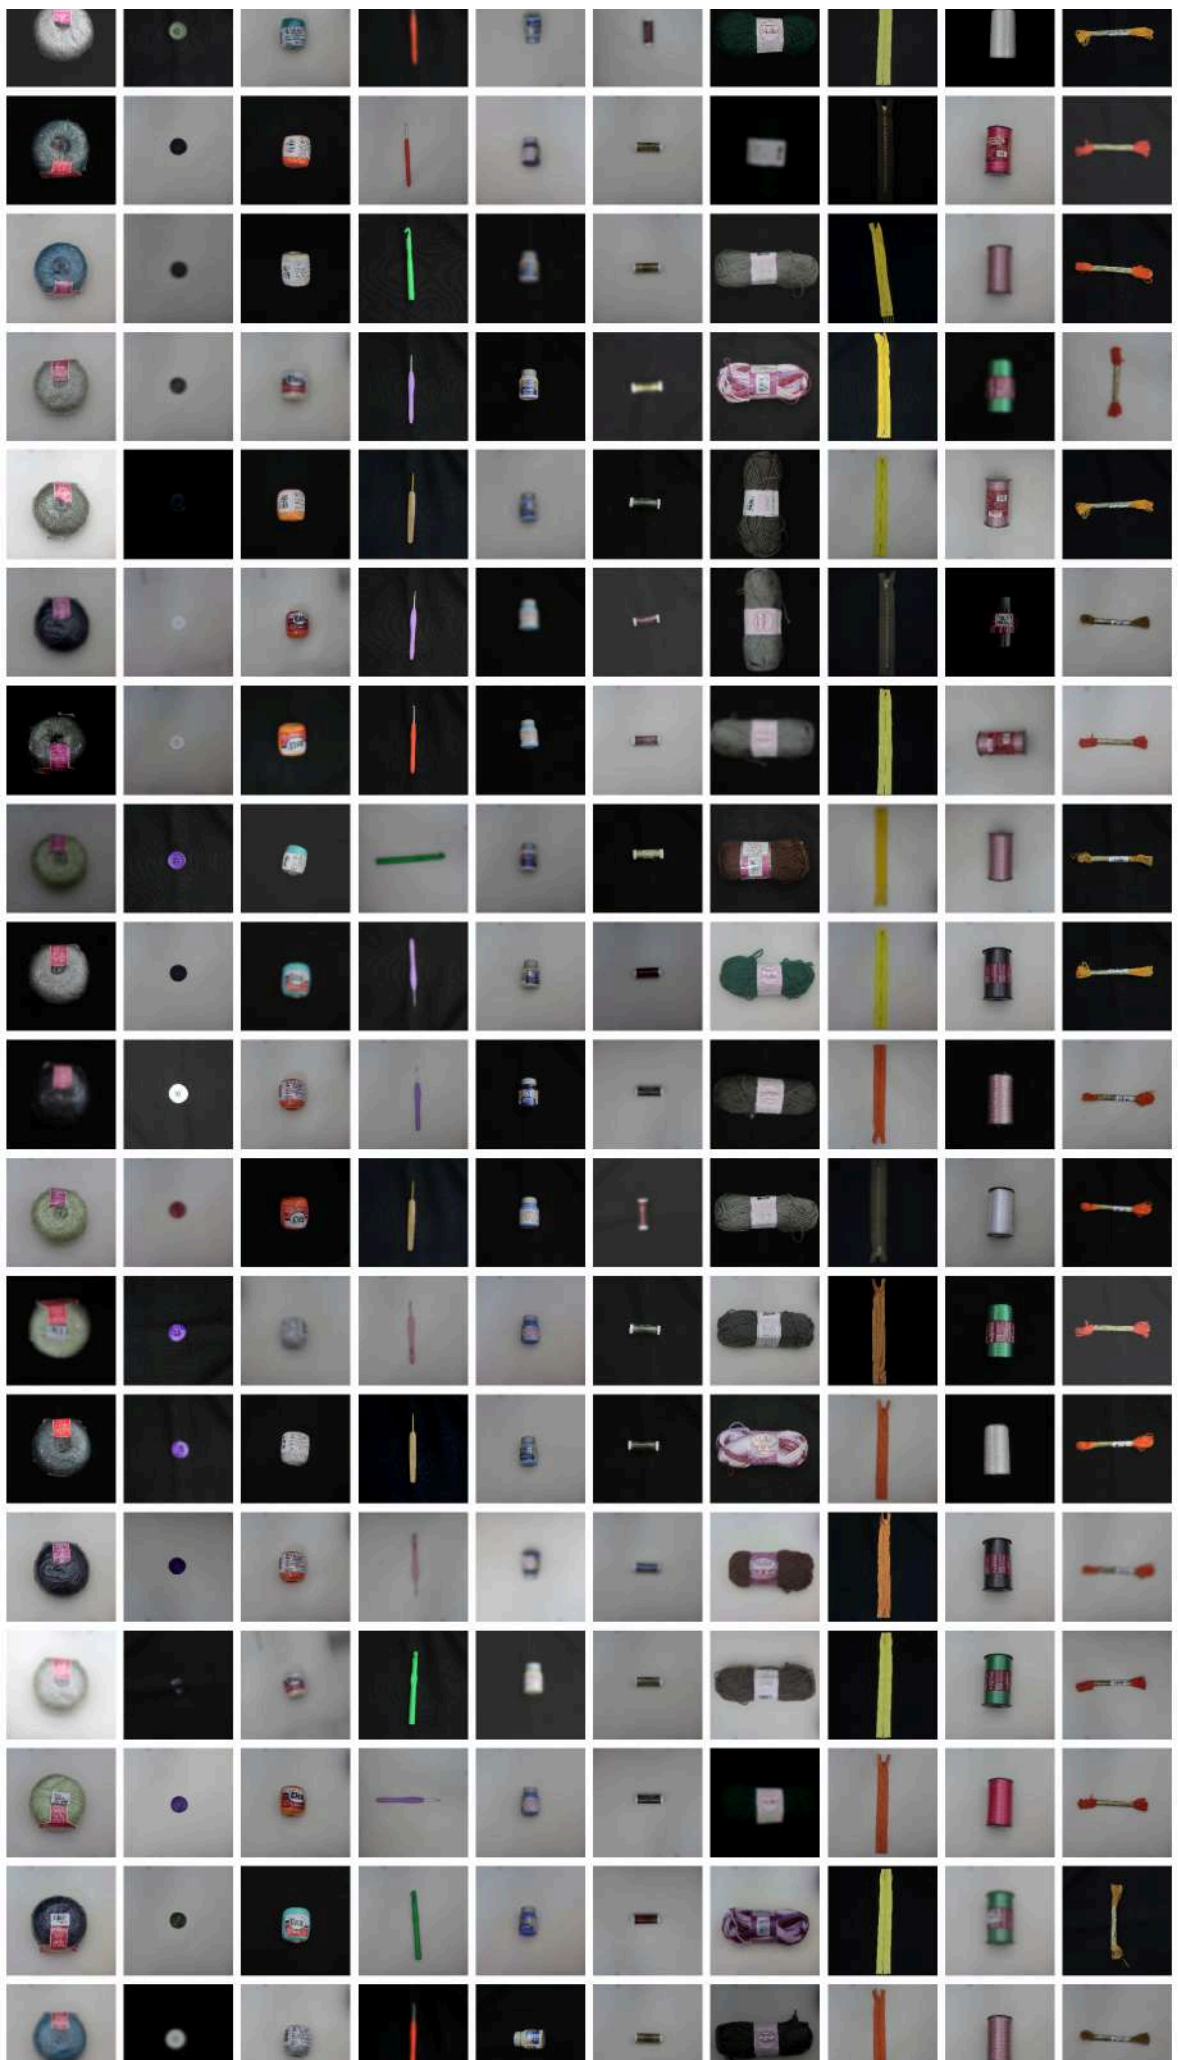
A segunda etapa envolve a aumentação do dataset, realizada pela função `dataset_augmentation` dentro da classe `TrimDataset`. Utilizando a biblioteca Albumentations, tal função aplica uma série de transformações a cada imagem do dataset, sendo elas: Uma dentre quatro transformações geométricas possíveis, ajuste aleatório do brilho e do contraste (50% de chance), blur gaussiano (50% de chance), ajuste da saturação da imagem (50% de chance).

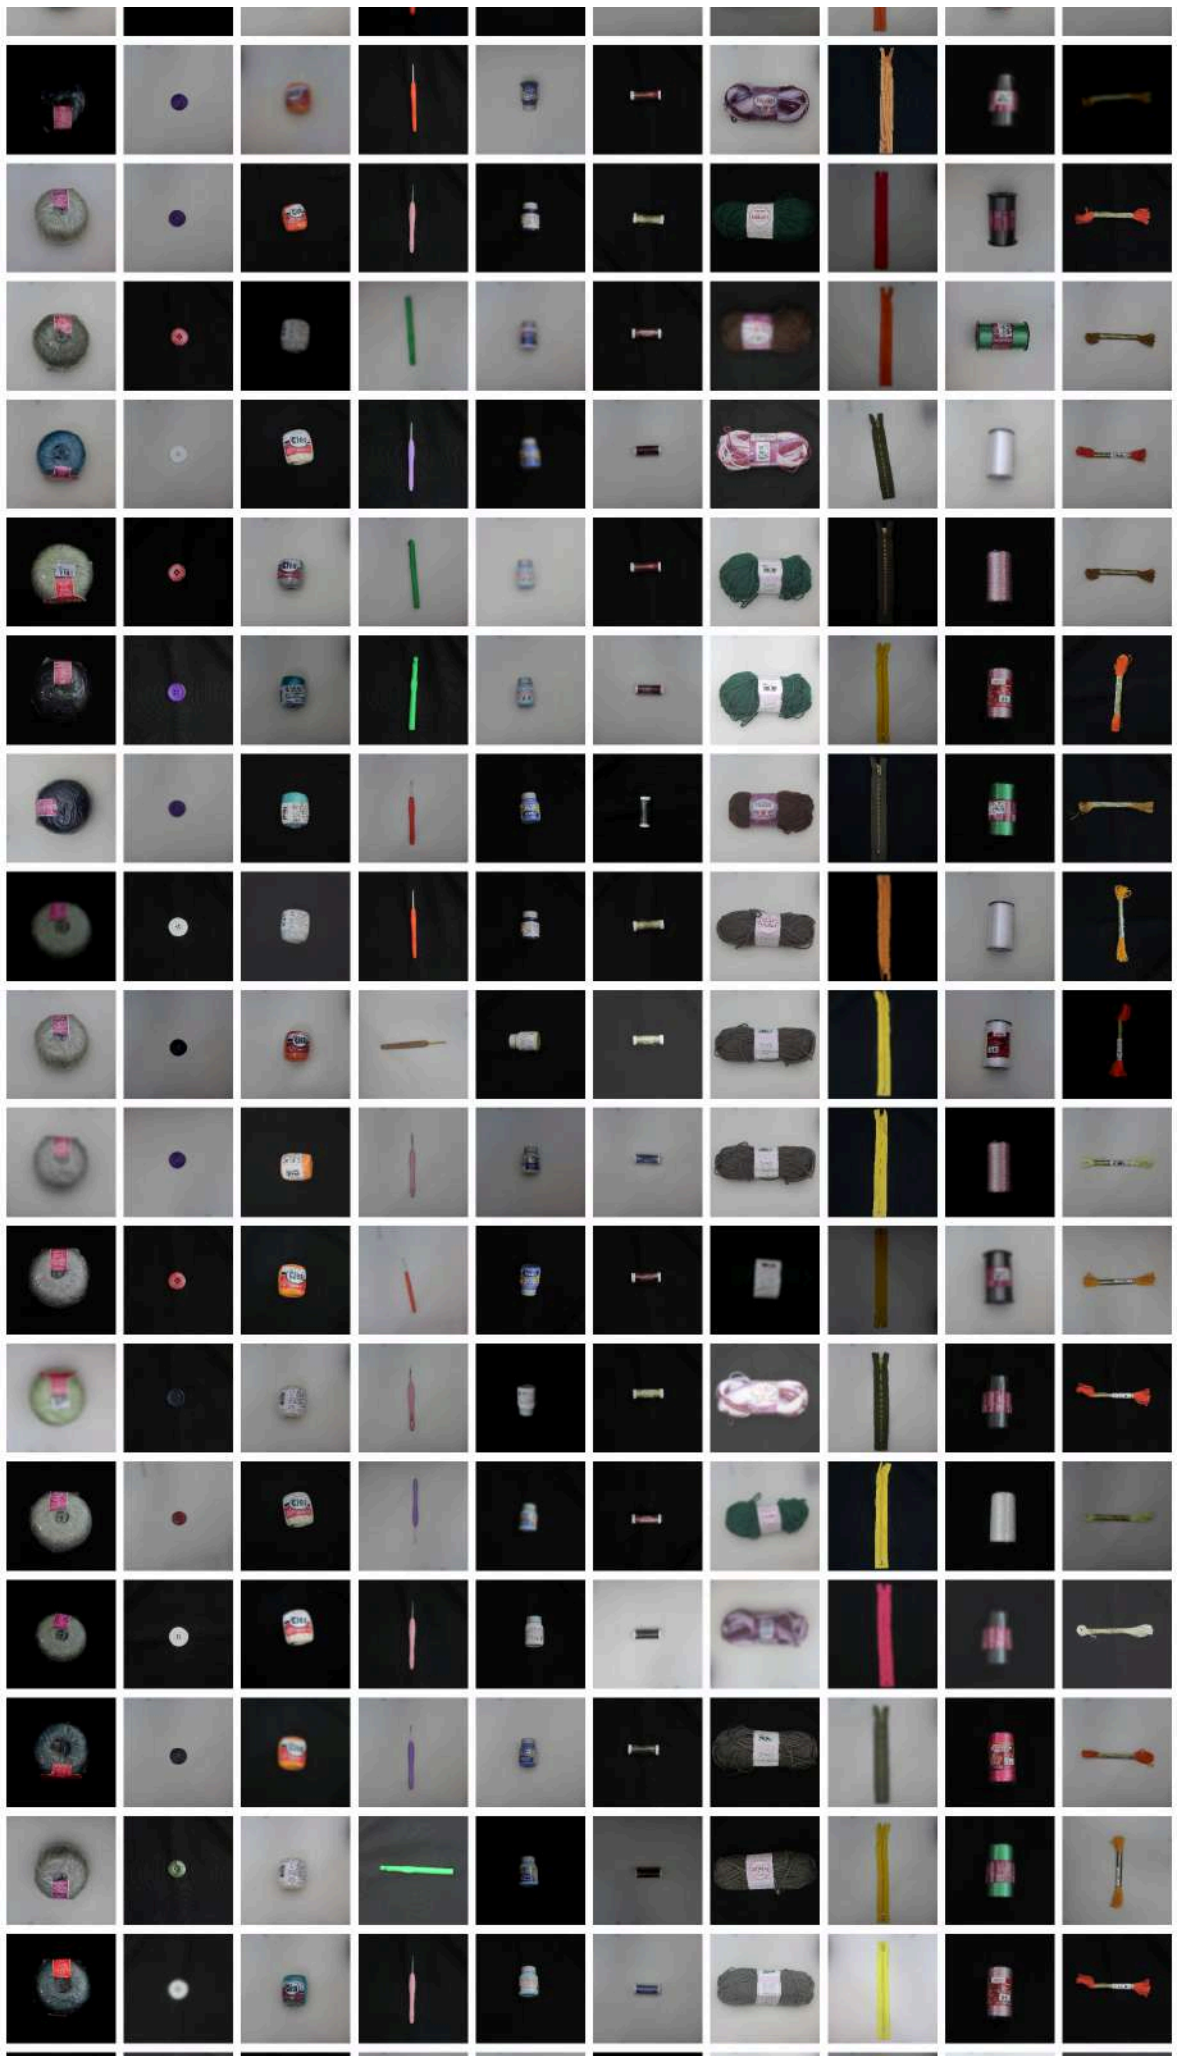
Ao fim dessa etapa, o dataset aumenta de 200 fotos originais para 600 fotos aumentadas, sendo que essa nova versão do dataset é salvo em uma nova pasta denominada `augmented_dataset`.

No pedaço de código abaixo, executamos o `dataset_augmentation` e plotamos todas as fotos que compõem o dataset até o momento.


```
In [4]: dataset.dataset_augmentation()  
plot_class_grid(dataset, 60, "nobbox")
```











A terceira e última etapa do pré-processamento é a normalização do histograma das imagens do dataset, realizada pela função `normalize_dataset` dentro da classe `TrimDataset`. Essa função aplica equalização do histograma para cada imagem do dataset, salvando as novas imagens em uma outra pasta denominada `normalized_dataset`.

```
In [5]: dataset.normalize_dataset()
```


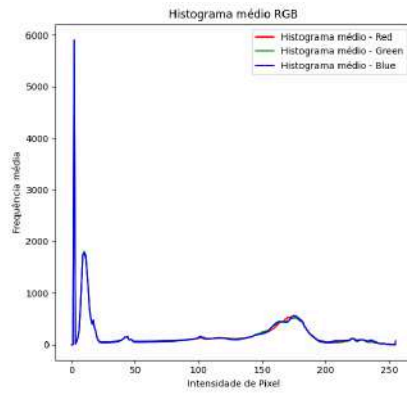

Após a etapa de normalização, é necessário recarregar o dataset, isto se dá pois as novas imagens normalizadas são salvas na pasta mas não são armazenadas no objeto `dataset`.

```
In [6]: new_dataset = TrimDataset('normalized_dataset')
```




Com o dataset normalizado, abaixo podemos ver o protótipo médio, histograma médio, e a variância do histograma de cada classe.

```
In [7]: generate_image_statistics(new_dataset.images, new_dataset.categories)
```




Protótipo médio amigurumi




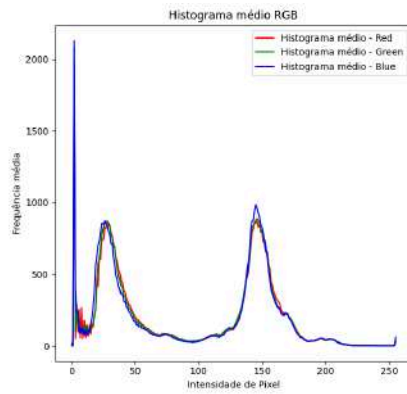
Protótipo médio button



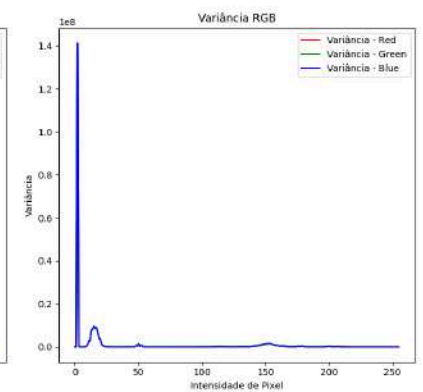

Protótipo médio clea




Protótipo médio crochet\_needle



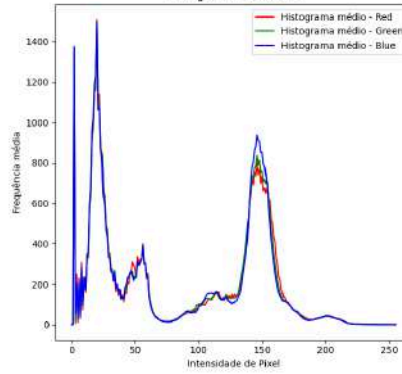
Protótipo médio paint




Protótipo médio spool




Histograma médio RGB




Variância RGB




Protótipo médio wool



Histograma médio RGB



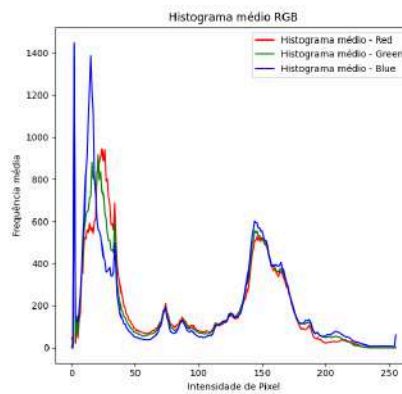
Variância RGB




Protótipo médio zipper



Histograma médio RGB




Variância RGB




Protótipo médio ribbon




Histograma médio RGB




Variância RGB




Protótipo médio skein



Histograma médio RGB




Variância RGB





No CVAT, além da marcação das bounding boxes, também realizamos a segmentação de todas as imagens, informação esta que foi carregada para as etapas de augmentation e normalization. Dessa forma, o dataset final contém o Ground Truth de todas as imagens, mesmo que apenas a bounding box seja utilizada para o treinamento do modelo na etapa posterior. O código abaixo gera o ground truth de 5 imagens aleatórias de cada classe.

```
In [8]: plot_class_grid(new_dataset, 5, "gt_mask")
```



## Parte 2: Treinamento do classificador utilizando CNN

Utilizando a biblioteca TensorFlow, treinamos um modelo classificador baseado em CNN.

```
In [9]: input_shape = (224,224,3) # Tamanho fixo das imagens no Dataset
num_classes = 10 # Numero de classes para classificar

model = create_cnn(input_shape, num_classes) # classificador a ser treinado
```

Antes de treinarmos o modelo, as imagens devem ser organizadas, separadas e pré-processadas para se adequarem aos formatos utilizados pela biblioteca TensorFlow. Primeiramente, organizamos todas as imagens, bbox e categorias na forma de um dicionário, então, dividimos os dados para cada etapa: 80% para treinamento, 10% para teste e 10% para validação.

```
In [10]: images = [data_to_dict(img.content, img.bbox, img.category_id) for img in new_da
img_train, img_test, img_validation = split_dataset(images)
```

Agora, separamos as informações nos conjuntos de treino teste e validação, isto é, para cada conjunto, separamos as informações das imagens, classes e bbox em vetores, pré-processando-os para serem utilizados pela biblioteca da TensorFlow. Além disso, as imagens devem ser convertidas do formato BGR do cv2 para o formato RGB, normalizado de 0 a 1, para serem utilizadas pelo modelo.

```

In [11]: from tensorflow.python.keras.utils.np_utils import to_categorical
import numpy as np

# Imagens de Treino
x_train = [img['image'] for img in img_train]
x_train = preprocess_images(x_train)

# Classes das imagens de treino
y_train_classes = np.array([img['class'] - 1 for img in img_train])
y_train_classes = to_categorical(y_train_classes, num_classes)

# Bboxes das imagens de treino
y_train_bbox = [img['bbox'] for img in img_train]
y_train_bbox = np.array(y_train_bbox) #/ [224.0, 224.0, 224.0, 224.0]

# Mesmo processo para imagens de validação e teste
x_val = [img['image'] for img in img_validation]
x_val = preprocess_images(x_val)
y_val_classes = np.array([img['class'] - 1 for img in img_validation])
y_val_classes = to_categorical(y_val_classes, num_classes)
y_val_bbox = [img['bbox'] for img in img_validation]
y_val_bbox = np.array(y_val_bbox) #/ [224.0, 224.0, 224.0, 224.0]

x_test = [img['image'] for img in img_test]
x_test = preprocess_images(x_test)
y_test_classes = np.array([img['class'] - 1 for img in img_test])
y_test_classes = to_categorical(y_test_classes, num_classes)
y_test_bbox = [img['bbox'] for img in img_test]
y_test_bbox = np.array(y_test_bbox) #/ [224.0, 224.0, 224.0, 224.0]

```

Agora finalmente podemos treinar o modelo. Os valores de `epoch` e `batch_size` foram escolhidos empiricamente por meio de diversas tentativas até que o modelo atingisse bons resultados nas etapas de teste e validação.

```

In [12]: history = model.fit(
    x_train,
    {'class_output': y_train_classes, 'bbox_output': y_train_bbox},
    epochs=50,
    batch_size=32,
    validation_data=(x_val, {'class_output': y_val_classes, 'bbox_output': y_val
    })

```

Epoch 1/50

**15/15** ————— **13s** 623ms/step - bbox\_output\_loss: 5604.3750 - bbox\_output\_mean\_squared\_error: 5604.3750 - class\_output\_accuracy: 0.1106 - class\_output\_loss: 7.3032 - loss: 5611.6782 - val\_bbox\_output\_loss: 2376.2495 - val\_bbox\_output\_mean\_squared\_error: 2365.2156 - val\_class\_output\_accuracy: 0.1167 - val\_class\_output\_loss: 8.1233 - val\_loss: 2373.3506

Epoch 2/50

**15/15** ————— **9s** 630ms/step - bbox\_output\_loss: 2168.9563 - bbox\_output\_mean\_squared\_error: 2168.9563 - class\_output\_accuracy: 0.0921 - class\_output\_loss: 8.9122 - loss: 2177.8687 - val\_bbox\_output\_loss: 1633.5156 - val\_bbox\_output\_mean\_squared\_error: 1630.2109 - val\_class\_output\_accuracy: 0.0000e+00 - val\_class\_output\_loss: 10.9990 - val\_loss: 1641.2677

Epoch 3/50

**15/15** ————— **11s** 708ms/step - bbox\_output\_loss: 1635.6621 - bbox\_output\_mean\_squared\_error: 1635.6621 - class\_output\_accuracy: 0.0953 - class\_output\_loss: 10.1614 - loss: 1645.8234 - val\_bbox\_output\_loss: 898.5214 - val\_bbox\_output\_mean\_squared\_error: 902.5738 - val\_class\_output\_accuracy: 0.0667 - val\_class\_output\_loss: 8.3130 - val\_loss: 910.9404

Epoch 4/50

**15/15** ————— **11s** 722ms/step - bbox\_output\_loss: 813.1842 - bbox\_output\_mean\_squared\_error: 813.1842 - class\_output\_accuracy: 0.1589 - class\_output\_loss: 8.5129 - loss: 821.6970 - val\_bbox\_output\_loss: 532.9447 - val\_bbox\_output\_mean\_squared\_error: 534.9633 - val\_class\_output\_accuracy: 0.3167 - val\_class\_output\_loss: 7.8876 - val\_loss: 542.9113

Epoch 5/50

**15/15** ————— **11s** 758ms/step - bbox\_output\_loss: 531.8262 - bbox\_output\_mean\_squared\_error: 531.8262 - class\_output\_accuracy: 0.1965 - class\_output\_loss: 6.8320 - loss: 538.6582 - val\_bbox\_output\_loss: 405.3309 - val\_bbox\_output\_mean\_squared\_error: 410.0552 - val\_class\_output\_accuracy: 0.1667 - val\_class\_output\_loss: 5.2925 - val\_loss: 415.4068

Epoch 6/50

**15/15** ————— **11s** 741ms/step - bbox\_output\_loss: 465.0404 - bbox\_output\_mean\_squared\_error: 465.0404 - class\_output\_accuracy: 0.2311 - class\_output\_loss: 4.6225 - loss: 469.6629 - val\_bbox\_output\_loss: 301.5461 - val\_bbox\_output\_mean\_squared\_error: 304.8027 - val\_class\_output\_accuracy: 0.3500 - val\_class\_output\_loss: 3.3167 - val\_loss: 308.1859

Epoch 7/50

**15/15** ————— **11s** 726ms/step - bbox\_output\_loss: 351.7140 - bbox\_output\_mean\_squared\_error: 351.7140 - class\_output\_accuracy: 0.3598 - class\_output\_loss: 2.8411 - loss: 354.5551 - val\_bbox\_output\_loss: 255.1482 - val\_bbox\_output\_mean\_squared\_error: 257.8760 - val\_class\_output\_accuracy: 0.4167 - val\_class\_output\_loss: 2.5585 - val\_loss: 260.4806

Epoch 8/50

**15/15** ————— **11s** 740ms/step - bbox\_output\_loss: 263.4869 - bbox\_output\_mean\_squared\_error: 263.4869 - class\_output\_accuracy: 0.4463 - class\_output\_loss: 2.5218 - loss: 266.0087 - val\_bbox\_output\_loss: 241.0575 - val\_bbox\_output\_mean\_squared\_error: 243.8037 - val\_class\_output\_accuracy: 0.4167 - val\_class\_output\_loss: 2.4652 - val\_loss: 246.3093


Epoch 9/50

**15/15** ————— **11s** 720ms/step - bbox\_output\_loss: 273.0204 - bbox\_output\_mean\_squared\_error: 273.0204 - class\_output\_accuracy: 0.5197 - class\_output\_loss: 2.1720 - loss: 275.1924 - val\_bbox\_output\_loss: 178.5638 - val\_bbox\_output\_mean\_squared\_error: 180.8592 - val\_class\_output\_accuracy: 0.4833 - val\_class\_output\_loss: 1.9661 - val\_loss: 182.8527


Epoch 10/50

**15/15** ————— **11s** 730ms/step - bbox\_output\_loss: 189.3588 - bbox\_output\_mean\_squared\_error: 189.3588 - class\_output\_accuracy: 0.5457 - class\_output\_loss: 1.8060 - loss: 191.1648 - val\_bbox\_output\_loss: 206.3238 - val\_bbox\_output\_mean\_squared\_error: 207.9528 - val\_class\_output\_accuracy: 0.5833 - val\_class\_output\_loss: 1.6983 - val\_loss: 209.6672


Epoch 11/50

**15/15**  **11s** 737ms/step - bbox\_output\_loss: 223.9917 - bbox\_output\_mean\_squared\_error: 223.9917 - class\_output\_accuracy: 0.6172 - class\_output\_loss: 1.6552 - loss: 225.6470 - val\_bbox\_output\_loss: 190.8862 - val\_bbox\_output\_mean\_squared\_error: 193.0817 - val\_class\_output\_accuracy: 0.5333 - val\_class\_output\_loss: 1.8020 - val\_loss: 194.8918


Epoch 12/50

**15/15**  **11s** 701ms/step - bbox\_output\_loss: 183.0847 - bbox\_output\_mean\_squared\_error: 183.0847 - class\_output\_accuracy: 0.5752 - class\_output\_loss: 2.1685 - loss: 185.2532 - val\_bbox\_output\_loss: 179.4358 - val\_bbox\_output\_mean\_squared\_error: 181.2851 - val\_class\_output\_accuracy: 0.5500 - val\_class\_output\_loss: 2.9245 - val\_loss: 184.2230


Epoch 13/50

**15/15**  **11s** 715ms/step - bbox\_output\_loss: 194.5088 - bbox\_output\_mean\_squared\_error: 194.5088 - class\_output\_accuracy: 0.6004 - class\_output\_loss: 2.1343 - loss: 196.6431 - val\_bbox\_output\_loss: 243.1518 - val\_bbox\_output\_mean\_squared\_error: 246.5459 - val\_class\_output\_accuracy: 0.6500 - val\_class\_output\_loss: 1.1962 - val\_loss: 247.7643


Epoch 14/50

**15/15**  **11s** 740ms/step - bbox\_output\_loss: 184.3775 - bbox\_output\_mean\_squared\_error: 184.3775 - class\_output\_accuracy: 0.7171 - class\_output\_loss: 1.4868 - loss: 185.8643 - val\_bbox\_output\_loss: 186.5235 - val\_bbox\_output\_mean\_squared\_error: 189.1734 - val\_class\_output\_accuracy: 0.7000 - val\_class\_output\_loss: 1.0034 - val\_loss: 190.2016


Epoch 15/50

**15/15**  **11s** 748ms/step - bbox\_output\_loss: 153.6552 - bbox\_output\_mean\_squared\_error: 153.6552 - class\_output\_accuracy: 0.8083 - class\_output\_loss: 0.8341 - loss: 154.4893 - val\_bbox\_output\_loss: 164.9910 - val\_bbox\_output\_mean\_squared\_error: 167.6374 - val\_class\_output\_accuracy: 0.6500 - val\_class\_output\_loss: 1.1355 - val\_loss: 168.7809


Epoch 16/50

**15/15**  **11s** 713ms/step - bbox\_output\_loss: 127.2298 - bbox\_output\_mean\_squared\_error: 127.2298 - class\_output\_accuracy: 0.7689 - class\_output\_loss: 0.9303 - loss: 128.1601 - val\_bbox\_output\_loss: 191.0435 - val\_bbox\_output\_mean\_squared\_error: 193.2492 - val\_class\_output\_accuracy: 0.5500 - val\_class\_output\_loss: 1.1285 - val\_loss: 194.3859


Epoch 17/50

**15/15**  **11s** 737ms/step - bbox\_output\_loss: 158.1259 - bbox\_output\_mean\_squared\_error: 158.1259 - class\_output\_accuracy: 0.7078 - class\_output\_loss: 1.1967 - loss: 159.3227 - val\_bbox\_output\_loss: 155.7765 - val\_bbox\_output\_mean\_squared\_error: 158.5487 - val\_class\_output\_accuracy: 0.6667 - val\_class\_output\_loss: 0.9405 - val\_loss: 159.5021


Epoch 18/50

**15/15**  **10s** 685ms/step - bbox\_output\_loss: 111.3629 - bbox\_output\_mean\_squared\_error: 111.3629 - class\_output\_accuracy: 0.8040 - class\_output\_loss: 0.9144 - loss: 112.2773 - val\_bbox\_output\_loss: 179.1065 - val\_bbox\_output\_mean\_squared\_error: 181.0321 - val\_class\_output\_accuracy: 0.5667 - val\_class\_output\_loss: 1.7631 - val\_loss: 182.8045

Epoch 19/50

**15/15**  **10s** 682ms/step - bbox\_output\_loss: 204.6862 - bbox\_output\_mean\_squared\_error: 204.6862 - class\_output\_accuracy: 0.7085 - class\_output\_loss: 1.0702 - loss: 205.7565 - val\_bbox\_output\_loss: 194.4577 - val\_bbox\_output\_mean\_squared\_error: 197.0762 - val\_class\_output\_accuracy: 0.7500 - val\_class\_output\_loss: 1.1219 - val\_loss: 198.2133

Epoch 20/50

**15/15**  **11s** 726ms/step - bbox\_output\_loss: 136.2652 - bbox\_output\_mean\_squared\_error: 136.2652 - class\_output\_accuracy: 0.7964 - class\_output\_loss: 1.0039 - loss: 137.2691 - val\_bbox\_output\_loss: 160.5117 - val\_bbox\_output\_mean\_squared\_error: 163.9299 - val\_class\_output\_accuracy: 0.6833 - val\_class\_output\_loss: 1.1420 - val\_loss: 165.0791

Epoch 21/50

**15/15** ————— **10s** 659ms/step - bbox\_output\_loss: 113.7127 - bbox\_output\_mean\_squared\_error: 113.7127 - class\_output\_accuracy: 0.8013 - class\_output\_loss: 0.8173 - loss: 114.5300 - val\_bbox\_output\_loss: 150.3305 - val\_bbox\_output\_mean\_squared\_error: 152.9409 - val\_class\_output\_accuracy: 0.6667 - val\_class\_output\_loss: 1.0492 - val\_loss: 154.0171

Epoch 22/50

**15/15** ————— **11s** 719ms/step - bbox\_output\_loss: 107.6781 - bbox\_output\_mean\_squared\_error: 107.6781 - class\_output\_accuracy: 0.8161 - class\_output\_loss: 0.7585 - loss: 108.4366 - val\_bbox\_output\_loss: 143.8892 - val\_bbox\_output\_mean\_squared\_error: 146.0589 - val\_class\_output\_accuracy: 0.7333 - val\_class\_output\_loss: 0.6720 - val\_loss: 146.7471

Epoch 23/50

**15/15** ————— **12s** 842ms/step - bbox\_output\_loss: 97.3185 - bbox\_output\_mean\_squared\_error: 97.3185 - class\_output\_accuracy: 0.8289 - class\_output\_loss: 0.6205 - loss: 97.9390 - val\_bbox\_output\_loss: 137.2159 - val\_bbox\_output\_mean\_squared\_error: 140.1492 - val\_class\_output\_accuracy: 0.8167 - val\_class\_output\_loss: 0.7031 - val\_loss: 140.8755

Epoch 24/50

**15/15** ————— **13s** 883ms/step - bbox\_output\_loss: 82.6470 - bbox\_output\_mean\_squared\_error: 82.6470 - class\_output\_accuracy: 0.8365 - class\_output\_loss: 0.5540 - loss: 83.2010 - val\_bbox\_output\_loss: 192.7413 - val\_bbox\_output\_mean\_squared\_error: 195.3733 - val\_class\_output\_accuracy: 0.6667 - val\_class\_output\_loss: 0.8959 - val\_loss: 196.2780

Epoch 25/50

**15/15** ————— **17s** 630ms/step - bbox\_output\_loss: 91.6059 - bbox\_output\_mean\_squared\_error: 91.6059 - class\_output\_accuracy: 0.8592 - class\_output\_loss: 0.5884 - loss: 92.1943 - val\_bbox\_output\_loss: 144.9444 - val\_bbox\_output\_mean\_squared\_error: 147.6810 - val\_class\_output\_accuracy: 0.8333 - val\_class\_output\_loss: 0.6114 - val\_loss: 148.3113

Epoch 26/50

**15/15** ————— **11s** 766ms/step - bbox\_output\_loss: 84.2213 - bbox\_output\_mean\_squared\_error: 84.2213 - class\_output\_accuracy: 0.8554 - class\_output\_loss: 0.5266 - loss: 84.7479 - val\_bbox\_output\_loss: 136.4364 - val\_bbox\_output\_mean\_squared\_error: 139.2730 - val\_class\_output\_accuracy: 0.6333 - val\_class\_output\_loss: 0.8544 - val\_loss: 140.1534

Epoch 27/50

**15/15** ————— **10s** 674ms/step - bbox\_output\_loss: 78.6004 - bbox\_output\_mean\_squared\_error: 78.6004 - class\_output\_accuracy: 0.8625 - class\_output\_loss: 0.4531 - loss: 79.0535 - val\_bbox\_output\_loss: 123.0003 - val\_bbox\_output\_mean\_squared\_error: 125.3670 - val\_class\_output\_accuracy: 0.7833 - val\_class\_output\_loss: 0.5670 - val\_loss: 125.9493

Epoch 28/50

**15/15** ————— **10s** 646ms/step - bbox\_output\_loss: 66.8144 - bbox\_output\_mean\_squared\_error: 66.8144 - class\_output\_accuracy: 0.8810 - class\_output\_loss: 0.4086 - loss: 67.2230 - val\_bbox\_output\_loss: 129.4022 - val\_bbox\_output\_mean\_squared\_error: 132.2429 - val\_class\_output\_accuracy: 0.6833 - val\_class\_output\_loss: 0.7564 - val\_loss: 133.0180

Epoch 29/50


**15/15** ————— **11s** 722ms/step - bbox\_output\_loss: 71.3705 - bbox\_output\_mean\_squared\_error: 71.3705 - class\_output\_accuracy: 0.8594 - class\_output\_loss: 0.4837 - loss: 71.8542 - val\_bbox\_output\_loss: 131.6530 - val\_bbox\_output\_mean\_squared\_error: 134.4069 - val\_class\_output\_accuracy: 0.8000 - val\_class\_output\_loss: 0.5524 - val\_loss: 134.9756

Epoch 30/50


**15/15** ————— **11s** 726ms/step - bbox\_output\_loss: 61.1270 - bbox\_output\_mean\_squared\_error: 61.1270 - class\_output\_accuracy: 0.8841 - class\_output\_loss: 0.3751 - loss: 61.5021 - val\_bbox\_output\_loss: 124.6608 - val\_bbox\_output\_mean\_squared\_error: 127.3000 - val\_class\_output\_accuracy: 0.7667 - val\_class\_output\_loss: 0.6522 - val\_loss: 127.9735




Epoch 31/50

**15/15**  **10s** 668ms/step - bbox\_output\_loss: 61.4751 - bbox\_output\_mean\_squared\_error: 61.4751 - class\_output\_accuracy: 0.8695 - class\_output\_loss: 0.3804 - loss: 61.8555 - val\_bbox\_output\_loss: 126.6533 - val\_bbox\_output\_mean\_squared\_error: 128.8551 - val\_class\_output\_accuracy: 0.8500 - val\_class\_output\_loss: 0.4633 - val\_loss: 129.3383


Epoch 32/50

**15/15**  **12s** 771ms/step - bbox\_output\_loss: 71.0085 - bbox\_output\_mean\_squared\_error: 71.0085 - class\_output\_accuracy: 0.9017 - class\_output\_loss: 0.3875 - loss: 71.3961 - val\_bbox\_output\_loss: 122.0856 - val\_bbox\_output\_mean\_squared\_error: 124.7403 - val\_class\_output\_accuracy: 0.8667 - val\_class\_output\_loss: 0.5218 - val\_loss: 125.2833


Epoch 33/50

**15/15**  **10s** 667ms/step - bbox\_output\_loss: 65.2165 - bbox\_output\_mean\_squared\_error: 65.2165 - class\_output\_accuracy: 0.8286 - class\_output\_loss: 0.5935 - loss: 65.8100 - val\_bbox\_output\_loss: 137.6354 - val\_bbox\_output\_mean\_squared\_error: 140.3071 - val\_class\_output\_accuracy: 0.7833 - val\_class\_output\_loss: 0.7318 - val\_loss: 141.0610


Epoch 34/50

**15/15**  **9s** 634ms/step - bbox\_output\_loss: 57.0305 - bbox\_output\_mean\_squared\_error: 57.0305 - class\_output\_accuracy: 0.8759 - class\_output\_loss: 0.5014 - loss: 57.5319 - val\_bbox\_output\_loss: 119.3592 - val\_bbox\_output\_mean\_squared\_error: 121.2533 - val\_class\_output\_accuracy: 0.8167 - val\_class\_output\_loss: 0.4253 - val\_loss: 121.6874


Epoch 35/50

**15/15**  **10s** 688ms/step - bbox\_output\_loss: 47.9389 - bbox\_output\_mean\_squared\_error: 47.9389 - class\_output\_accuracy: 0.8876 - class\_output\_loss: 0.3724 - loss: 48.3113 - val\_bbox\_output\_loss: 115.1462 - val\_bbox\_output\_mean\_squared\_error: 117.5882 - val\_class\_output\_accuracy: 0.7667 - val\_class\_output\_loss: 0.6259 - val\_loss: 118.2352


Epoch 36/50

**15/15**  **10s** 657ms/step - bbox\_output\_loss: 47.0359 - bbox\_output\_mean\_squared\_error: 47.0359 - class\_output\_accuracy: 0.9210 - class\_output\_loss: 0.2932 - loss: 47.3291 - val\_bbox\_output\_loss: 126.2548 - val\_bbox\_output\_mean\_squared\_error: 129.2766 - val\_class\_output\_accuracy: 0.8667 - val\_class\_output\_loss: 0.4095 - val\_loss: 129.7007


Epoch 37/50

**15/15**  **10s** 666ms/step - bbox\_output\_loss: 49.8117 - bbox\_output\_mean\_squared\_error: 49.8117 - class\_output\_accuracy: 0.9023 - class\_output\_loss: 0.3248 - loss: 50.1365 - val\_bbox\_output\_loss: 129.5103 - val\_bbox\_output\_mean\_squared\_error: 130.8223 - val\_class\_output\_accuracy: 0.8333 - val\_class\_output\_loss: 0.4104 - val\_loss: 131.2437


Epoch 38/50

**15/15**  **10s** 664ms/step - bbox\_output\_loss: 53.2712 - bbox\_output\_mean\_squared\_error: 53.2712 - class\_output\_accuracy: 0.9126 - class\_output\_loss: 0.3336 - loss: 53.6048 - val\_bbox\_output\_loss: 126.2213 - val\_bbox\_output\_mean\_squared\_error: 128.2318 - val\_class\_output\_accuracy: 0.8500 - val\_class\_output\_loss: 0.4025 - val\_loss: 128.6460


Epoch 39/50

**15/15**  **9s** 631ms/step - bbox\_output\_loss: 45.0408 - bbox\_output\_mean\_squared\_error: 45.0408 - class\_output\_accuracy: 0.9043 - class\_output\_loss: 0.2797 - loss: 45.3205 - val\_bbox\_output\_loss: 131.4957 - val\_bbox\_output\_mean\_squared\_error: 133.4660 - val\_class\_output\_accuracy: 0.7833 - val\_class\_output\_loss: 0.5096 - val\_loss: 133.9927


Epoch 40/50

**15/15**  **10s** 648ms/step - bbox\_output\_loss: 41.1916 - bbox\_output\_mean\_squared\_error: 41.1916 - class\_output\_accuracy: 0.9026 - class\_output\_loss: 0.2994 - loss: 41.4910 - val\_bbox\_output\_loss: 110.4849 - val\_bbox\_output\_mean\_squared\_error: 112.6686 - val\_class\_output\_accuracy: 0.8667 - val\_class\_output\_loss: 0.3489 - val\_loss: 113.0276


Epoch 41/50

**15/15**  **14s** 939ms/step - bbox\_output\_loss: 36.2059 - bbox\_output\_mean\_squared\_error: 36.2059 - class\_output\_accuracy: 0.9165 - class\_output\_loss: 0.2410 - loss: 36.4470 - val\_bbox\_output\_loss: 134.6227 - val\_bbox\_output\_mean\_squared\_error: 136.9628 - val\_class\_output\_accuracy: 0.8667 - val\_class\_output\_loss: 0.4925 - val\_loss: 137.4757


Epoch 42/50

**15/15**  **11s** 738ms/step - bbox\_output\_loss: 31.8737 - bbox\_output\_mean\_squared\_error: 31.8737 - class\_output\_accuracy: 0.8873 - class\_output\_loss: 0.3390 - loss: 32.2127 - val\_bbox\_output\_loss: 119.5234 - val\_bbox\_output\_mean\_squared\_error: 121.7450 - val\_class\_output\_accuracy: 0.8167 - val\_class\_output\_loss: 0.4280 - val\_loss: 122.1825


Epoch 43/50

**15/15**  **15s** 1s/step - bbox\_output\_loss: 34.3518 - bbox\_output\_mean\_squared\_error: 34.3518 - class\_output\_accuracy: 0.9178 - class\_output\_loss: 0.2701 - loss: 34.6219 - val\_bbox\_output\_loss: 102.2769 - val\_bbox\_output\_mean\_squared\_error: 104.0714 - val\_class\_output\_accuracy: 0.8500 - val\_class\_output\_loss: 0.4231 - val\_loss: 104.5116


Epoch 44/50

**15/15**  **24s** 1s/step - bbox\_output\_loss: 37.4344 - bbox\_output\_mean\_squared\_error: 37.4344 - class\_output\_accuracy: 0.9182 - class\_output\_loss: 0.2554 - loss: 37.6897 - val\_bbox\_output\_loss: 120.8413 - val\_bbox\_output\_mean\_squared\_error: 122.6126 - val\_class\_output\_accuracy: 0.9000 - val\_class\_output\_loss: 0.2866 - val\_loss: 122.9072


Epoch 45/50

**15/15**  **11s** 697ms/step - bbox\_output\_loss: 31.4869 - bbox\_output\_mean\_squared\_error: 31.4869 - class\_output\_accuracy: 0.9344 - class\_output\_loss: 0.2544 - loss: 31.7413 - val\_bbox\_output\_loss: 106.8804 - val\_bbox\_output\_mean\_squared\_error: 108.7897 - val\_class\_output\_accuracy: 0.8667 - val\_class\_output\_loss: 0.3304 - val\_loss: 109.1345


Epoch 46/50

**15/15**  **10s** 693ms/step - bbox\_output\_loss: 26.7571 - bbox\_output\_mean\_squared\_error: 26.7571 - class\_output\_accuracy: 0.9322 - class\_output\_loss: 0.2076 - loss: 26.9647 - val\_bbox\_output\_loss: 105.2763 - val\_bbox\_output\_mean\_squared\_error: 107.2338 - val\_class\_output\_accuracy: 0.8667 - val\_class\_output\_loss: 0.3707 - val\_loss: 107.6166


Epoch 47/50

**15/15**  **10s** 693ms/step - bbox\_output\_loss: 27.6417 - bbox\_output\_mean\_squared\_error: 27.6417 - class\_output\_accuracy: 0.9306 - class\_output\_loss: 0.1950 - loss: 27.8367 - val\_bbox\_output\_loss: 106.5825 - val\_bbox\_output\_mean\_squared\_error: 108.1990 - val\_class\_output\_accuracy: 0.8667 - val\_class\_output\_loss: 0.2766 - val\_loss: 108.4847


Epoch 48/50

**15/15**  **10s** 693ms/step - bbox\_output\_loss: 33.9391 - bbox\_output\_mean\_squared\_error: 33.9391 - class\_output\_accuracy: 0.9405 - class\_output\_loss: 0.1934 - loss: 34.1326 - val\_bbox\_output\_loss: 119.1563 - val\_bbox\_output\_mean\_squared\_error: 120.8843 - val\_class\_output\_accuracy: 0.8833 - val\_class\_output\_loss: 0.2576 - val\_loss: 121.1539

Epoch 49/50

**15/15**  **10s** 678ms/step - bbox\_output\_loss: 41.6245 - bbox\_output\_mean\_squared\_error: 41.6245 - class\_output\_accuracy: 0.9548 - class\_output\_loss: 0.1726 - loss: 41.7972 - val\_bbox\_output\_loss: 109.0860 - val\_bbox\_output\_mean\_squared\_error: 110.8908 - val\_class\_output\_accuracy: 0.8667 - val\_class\_output\_loss: 0.3463 - val\_loss: 111.2523

Epoch 50/50

**15/15**  **10s** 693ms/step - bbox\_output\_loss: 41.5185 - bbox\_output\_mean\_squared\_error: 41.5185 - class\_output\_accuracy: 0.9379 - class\_output\_loss: 0.2125 - loss: 41.7310 - val\_bbox\_output\_loss: 106.8088 - val\_bbox\_output\_mean\_squared\_error: 108.2408 - val\_class\_output\_accuracy: 0.8833 - val\_class\_output\_loss: 0.3174 - val\_loss: 108.5692

Agora, avaliamos o modelo com imagens não vistas na etapa de treinamento

```
In [13]: model.evaluate(  
    x_test,  
    {'bbox_output': y_test_bbox, 'class_output': y_test_classes}  
)
```

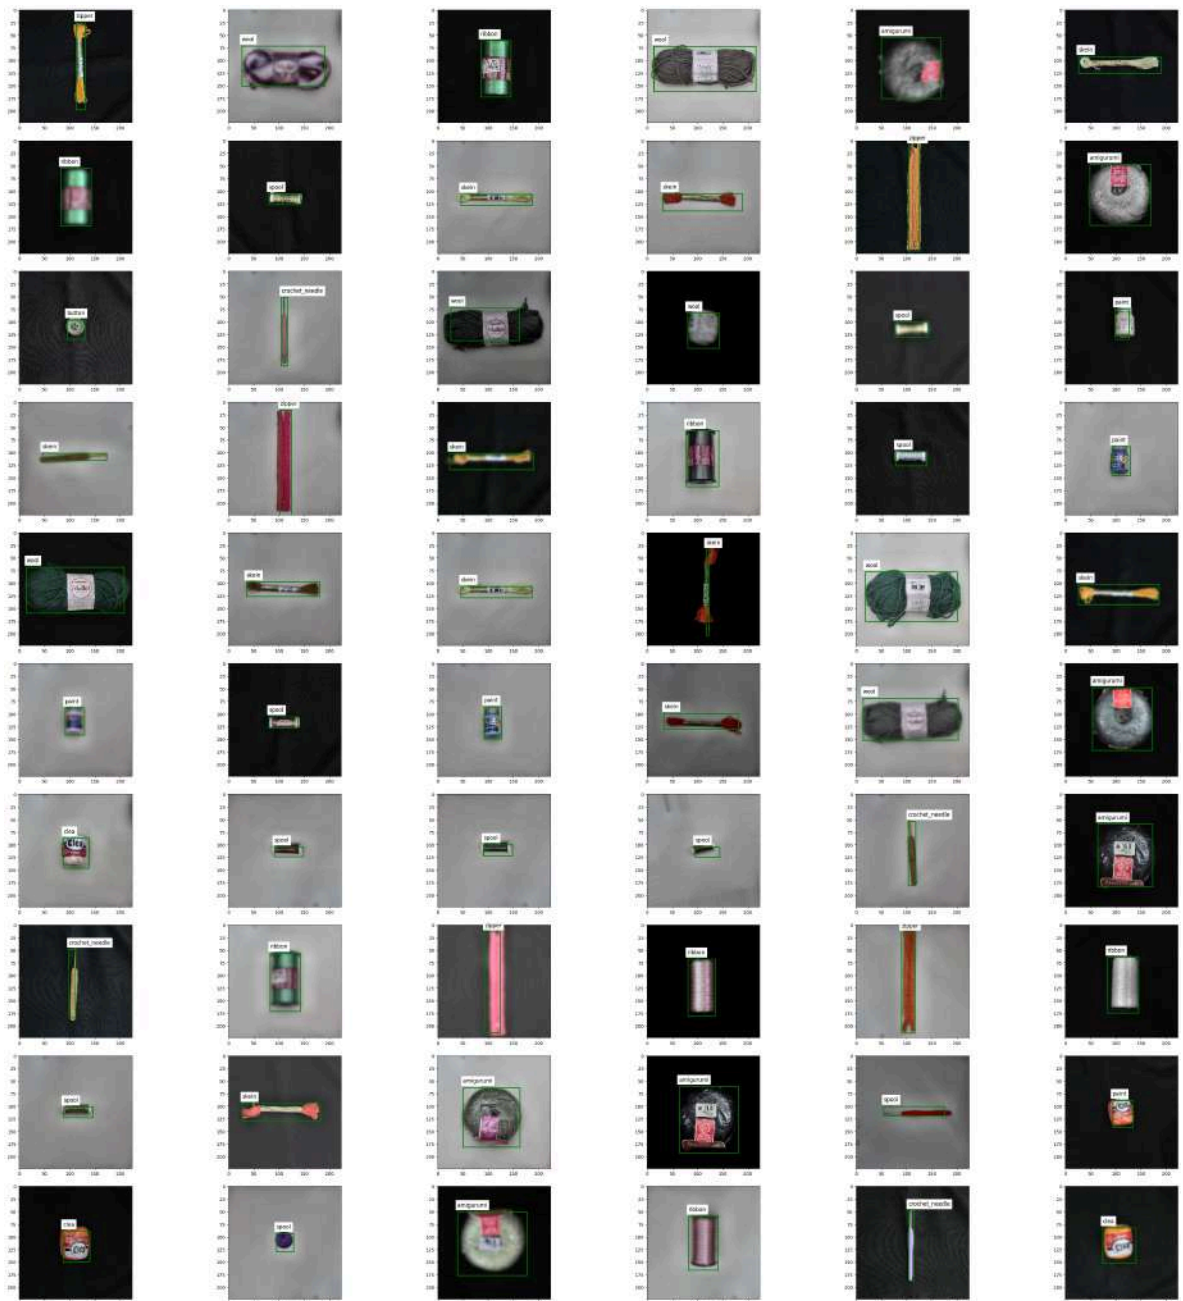
2/2 ————— 1s 191ms/step - bbox\_output\_loss: 76.7964 - bbox\_output\_mean\_squared\_error: 78.0613 - class\_output\_accuracy: 0.9021 - class\_output\_loss: 0.3149 - loss: 78.3751

```
Out[13]: [69.52832794189453,  
    0.3228673040866852,  
    67.3097152709961,  
    69.2070541381836,  
    0.8999999761581421]
```

Adicionalmente, podemos utilizar esses mesmos dados para visualizar a eficiência do modelo treinado.

```
In [14]: import matplotlib.pyplot as plt  
import matplotlib.patches as patches  
  
predicted_classes, predicted_bboxes = model.predict(x_test)  
predicted_classes = np.argmax(predicted_classes, axis=1) + 1  
  
true_classes = np.argmax(y_test_classes, axis=1) + 1  
class_labels = new_dataset.categories  
  
fig, axes = plt.subplots(10, 6, figsize=(40, 40))  
for i, ax in enumerate(axes.flat):  
    image = x_test[i]  
  
    bbox = predicted_bboxes[i]  
    predicted_class = predicted_classes[i]  
  
    ax.imshow(image)  
    rect = patches.Rectangle((bbox[0], bbox[1]), bbox[2], bbox[3],  
                             linewidth=2, edgecolor='g', facecolor='none')  
    ax.add_patch(rect)  
    ax.text(bbox[0], bbox[1] - 10.0, f'{class_labels[predicted_class]}', fontsize=12)  
  
plt.tight_layout()  
plt.show()
```

2/2 ————— 1s 248ms/step



Abaixo estão as métricas acurácia média e F1-Score, além da matriz de confusão do modelo.

```
In [15]: from sklearn.metrics import accuracy_score, f1_score, confusion_matrix, ConfusionMatrixDisplay

accuracy = accuracy_score(true_classes, predicted_classes)
print(f"Acuracia: {accuracy:.4f}")

f1 = f1_score(true_classes, predicted_classes, average='weighted')
print(f"F1-Score: {f1:.4f}")

cm = confusion_matrix(true_classes, predicted_classes)


fig, ax = plt.subplots(figsize=(5, 5))

categorias = list(new_dataset.categories.values())
ConfusionMatrixDisplay.from_predictions(
    true_classes, predicted_classes, display_labels=categorias, xticks_rotation=
    ax=ax, colorbar=True, cmap="plasma")
```

```
plt.show()
```

Acuracia: 0.9000

F1-Score: 0.8941



Se os resultados são satisfatórios, salvamos o modelo na pasta models.

```
In [16]: import os
```

```
model.save(os.path.join('models', 'cnn_bounding_box_model.keras'))
```